

Farfalle: a Domain-Specific Language and Toolsuite for Processor Microarchitectures

Alessandro de Gennaro

School of Electrical and
Electronic Engineering
Newcastle University

Newcastle upon Tyne, United Kingdom
Email: a.de-gennaro@ncl.ac.uk

Paulius Stankaitis

School of Computing Science
Newcastle University

Newcastle upon Tyne, United Kingdom
Email: paulius.stankaitis@ncl.ac.uk

Andrey Mokhov

School of Electrical and
Electronic Engineering
Newcastle University

Newcastle upon Tyne, United Kingdom
Email: andrey.mokhov@ncl.ac.uk

Abstract—Adopting a systematic approach for the design of processor instruction sets is instrumental for tackling the complexity of modern processors.

We present a tool-chain which follows the designer through all the design phases, from the specification of instructions to the hardware synthesis of a microcontroller. The flow is also meant to facilitate the understanding of the Instruction Set Architecture (ISA) reference manuals. Often, these documents are semi-formal, hard to read and fully understand, especially when dealing with highly concurrent implementations. We believe that designers will also benefit from a visual graph-based model, automatically derived from the ISA specification, and customisable to fit different needs. Some of the tools have already been developed and tested on ARMv6-M architecture, others yet need to be fully implemented. The design flow will be integrated into the Workcraft framework.

This paper presents a new domain-specific language for ISA specification and demonstrates how it can be used for the specification of ARMv6-M microcontroller. We also compare the presented approach with the others available in the literature.

I. INTRODUCTION

Technology progress allows industries to integrate always more transistors over the same amount of area, following Moore’s law. In turn, design complexity of these systems progressively increases. This led research to be focused on raising the level of abstraction of the languages used at the early-stages of design. This work presents a design flow for an easier specification, visualisation, simulation, customisation and hardware synthesis of instruction sets.

The consultation of an ISA reference manual (i.e. ARMv6-M [11]) can be difficult and tedious. Anthony Fox, in his attempt to describe a model of the ARMv7 instruction set, argues: “official reference manuals are large, stretching to many hundreds of pages - one can easily overlook subtle details or become bogged down with “uninteresting” background information” ([6], Section 1). And yet: “official descriptions are semi-formal (ambiguous)” ([6], Section 1).

In the light of the above, a simple and formal way to specify and, more specifically, visualise instruction sets is needed. A visual graph-based model can help designers for a quicker comprehension of the processor. ISAs in fact provide a software level description of the hardware itself. We use Conditional Partial Order Graph (CPOG) [1][2] as the visualisation

model. They can efficiently represent concurrent and sequential behaviours, and already come with a tested tool-kit for the customisation, encoding and hardware synthesis [3][4][5].

We propose a new domain-specific language for ISA specification. The current implementation is embedded in *Haskell* [12]. This is a high-level functional language that provides static type checking and inference capabilities, as well as useful predefined constructs and classes that we rely on in our implementation (i.e. the Monad class).

The article is organised as follows: Section II describes the framework we have developed, and how this can be used to interact with all the steps of the design flow. Section III presents a case-study: the ARMv6-M ISA and how it can be converted into a more comprehensible CPOG representation. And finally, Section IV summarises the achieved results, outlining the future research directions.

A. Related work

An attempt to use partial orders to describe the instructions of a complex hardware structure can be found in [10]. The author uses Conditional Partial Order Graphs to visually describe the instructions of the Intel 8051 Microcontroller. Yet, he used them for building an asynchronous controller and managing the internal execution flow of the datapath elements. The manual construction of partial orders is an error-prone process. Connecting all the operations, taking into account their dependencies and order, is a complex task. This further inspired our research, and led us to bridge this gap with an automated approach.

Regarding the language we chose to adopt, it is not difficult to find other cases where functional languages have been used as a starting point for ISA specification. In this direction, [8] provides an interesting attempt to build an infrastructure for instruction set development. The authors developed the concepts of *state* and *transformations*. The former represents the current state of the machine, which evolves over time according to the transformations (instructions). *F. Yuan and K. I. Eder* instead, created a formal and hierarchical model, which can be refined for fitting the needs of a particular ISA. This model (characterised in [7]) is composed by 4 abstract layers. Each of these layers describes a particular

aspect of the instruction set. The deeper the designer explores this hierarchical representation, the more refined will be the final system.

Another example can be found in [6]. The authors, in there, built a framework which can be extended to tailor instruction sets. They use the monadic programming style, based on three basic operations: *return*, *bind* and *parallel*. These are meant to mimic the flow of an entity (the hardware system), which is always returned as a result of the execution of an instruction. Instructions can be bound together either sequentially or in parallel. Here, the case study used is the ARMv7, a widespread instruction set embedded by the Cortex-A8 processor. An additional example of how an ISA might be specified inside a tool is present in [9]. Both the modules and the instructions here are introduced via a xml-based language, fairly readable but not very flexible.

In this work we aim at developing a domain-specific language based on Haskell. This will be adopted for the representation of the ARMv6-M, used as the case study. The target representation abstracts some details such as the bit-length of the registers or immediate values. The goal is the integration of the framework presented in Section II-A into a systematic design flow (Section II-B), and the simplification of the understanding of processor instruction sets with a graph-based model (Section III-C).

II. A DOMAIN-SPECIFIC LANGUAGE: FARFALLE

The language used to describe the framework is Haskell [12]. This is a functional language with high abstraction capabilities. We followed the path, already explored in [6], implementing the *monad class* to build the infrastructure of the domain-specific language (DSL).

In Farfalle, the processor is seen as an entity composed by a combination of an internal and external storage. Respectively, the former is represented by a set of registers, while the latter by a memory that contains both instructions and data; in compliance with the von Neumann architecture.

The processor executes instructions, that modify the state of the modelled machine.

$$P(Regs, Mem) \Rightarrow P'(Regs, Mem)$$

Every time the processor executes an instruction indeed, the state of the system evolves into a new state P' different from the previous one. Where either the registers and the memory might be different. This approach is relatively similar to the processor description in [8]. Therefore, we are confident that the high degree of abstraction obtained will allow the bidirectional connection of these two representations.

This section is organised as follow: Section II-A describes the structure of the framework. A characterisation of the design flow derived by this approach is presented in Section II-B, while a brief description of the verification side, related to Event-B interpretation, is discussed in Section II-C.

A. Structure of the framework

What marks our approach is the high degree of flexibility and abstraction that can tailor to every processor. The framework we present in this document is composed by different layers (modelled in various and interconnected modules). Each of them plays an important role in the structure.

The module in charge of modelling the low-level hardware functions is named *basics*. Positioned at very bottom layer, defines the ways through which the memory elements can be accessed, the behaviour of the data-path components, as well as the monad basic functions: *return* and *bind*.

Furthermore, the types used in the structure are defined. The registers and the memory are internally implemented with two maps of integer values. The former can be accessed either via the name of a special purpose register, or using the pattern $R\ Int$ (i.e. $R\ 5$, to point out the fifth location in the register file). Memory is accessed via an integer type address instead. Immediate values are also internally seen as integers. This approach is not accurate. Some of the features – the bit length of the registers, or the endianness of the data – are lost. Nonetheless, the high level of abstraction of the framework allows the interpretation of the model from several points of view (see Section II-B), and is eventually enough for the control unit synthesis.

Finally, *ComputationType* is defined. This is a flexible type in input to the arithmetic logic unit (ALU), which can be modified according to the ISA needs. As far as this unit is concerned in fact, different combinations of inputs can be given to it: one single register for an increment/decrement operation, either two registers, or one register and an immediate for an ADD/SUB/MUL operation, and so on. Therefore, *ComputationType* is not a fixed type and can be extended for fitting different ALU architectures.

Microprogram is the interface based on the low-level module just described. In this, just a few operations, generically present in a general purpose processor, are implemented. The function for incrementing the program counter and fetching the next instruction, the pop and push instructions, the interface methods for accessing the registers and memory (internally modelled in the module *basics*). In addition to the set of internal registers usually contained in a generic processor, this module defines three registers more: the program counter, the instruction register and the stack pointer.

ISA_model is the module where the user is supposed to introduce the custom instruction set architecture. Based on the bottom layers, this module can be customised adding other special purpose registers and extending ALU types. Most importantly, this will contain all the instructions one wants to include in a custom ISA. This will act as an interface for the software-level simulation and will be taken into account for the conversion towards other models and languages (Section II-B).

Main is the module meant for the high-level software simulation, in here all the instructions can be executed over the processor modelled. This might be useful for having a first proof of a consistent architecture. Indeed, the possible

implementation errors can be captured analysing the behaviour of the whole system.

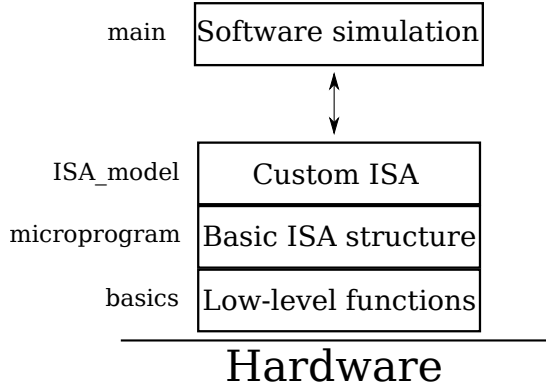


Fig. 1: Functional structure of the framework.

In Figure 1, the structure of the framework is depicted, pinpointing on the interconnections between the modules described. The structure is rather simple and small, and can be potentially extended with new procedures and types, making this approach flexible and scalable.

In Section III we will show how this framework has been implemented and extended to model the ARMv6-M.

B. Functionalities & interpretations

Our main goal with this article is to open the path towards an easier specification and visualisation of instruction sets. Nonetheless, quite a few other interpretations can be given to this work. The domain-specific language indeed, once used to model a processor, can be potentially translated into other languages. In turn, they can be progressively employed with the advantages and tools they come with. It is worth mentioning that some of the functionalities discussed in this section have not been implemented yet. Though, they will be part of our future work (see Section IV-A).

The ISA specification might be converted into Event-B language [13]. This is a formal technique used to analyse the model at the system level. It comes with some theorem provers which are able to verify the consistency of the system, proving that there are no inconsistencies between instructions. This is an interesting feature itself, which could be particularly useful when applied to instruction sets. It falls within the side of the formal *verification* (see Section II-C for a detailed description).

Visualisation is another interpretation which can be given to this framework. Even though Haskell provides readable and simple constructs, instruction set specifications are intrinsically difficult to read and understand. Capturing all the dependencies between the micro-operations within instructions is complex and error-prone. Reading and understanding the whole reference manual can take a long time. In the light of the above, we decided to use the conditional partial order graph model to represent the behaviour of the ISA. A partial order indeed, can be extracted by the formal specification of each instruction,

providing an easier environment for a better understanding by designers.

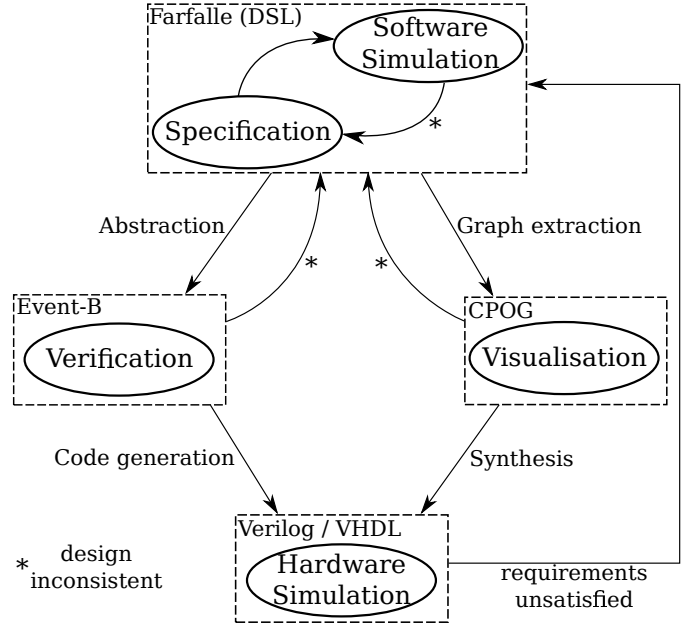


Fig. 2: Design flow supported by Farfalle.

This allows to employ the tool-chain around the conditional partial order graphs, enabling the visual customisation of the dependencies of the instructions in *Workcraft*, and the usage of CPOG-based *synthesis* tools. A fairly accurate measure of the size of the final microcontroller for the processor control unit can be automatically derived therefore (with a deep *hardware simulation*), allowing the designer to come back at the specification phase if some of the requirements (area, power consumption) are not met.

The whole design-flow is illustrated in Figure 2.

C. Verification with Event-B

Design validation process is instrumental for any kind of systems. A typical system design process requires numerous minor optimization modifications. The system needs to be re-validated after each refinement step, making this phase not particularly cost and time effective. Applying formal techniques is a complex task, which requires a deep mathematical understanding and rigorous analysis of the system. Furthermore, training domain specific engineers to mathematical techniques can be very expensive. Therefore, our approach will include methods and tool support, which intends to hide mathematical complexity by automating validation process.

Event-B [15], originated from well known formalism called B method, is one of the most used specification languages in safety critical software development area. Over the years, Event-B proved to be successful mainly due to a few large academic research projects. Nowadays, this is commonly used in industry. Excellent tool support and proof driven process [16] were key aspects which led us to choose this particular method.

The formalism is based on First Order Logic and Zermelo-Fraenkel set theory. This is a high-level and expressive language. System specification is the initial step in formal validation and it starts with describing requirements for the design. System requirements should completely cover functionality and environment of the design. Next stage is represented by the system development. Here, the model is developed: starting with the most abstract model, this is gradually refined broadening the level of granularity.

Formal verification is another major branch of formal modelling. After the system is specified and abstracted, its conditions need to be rigorously proved. Model is considered correct if and only if all conditions, also known as *proof obligations*, are satisfied. Beforehand, designer is required to define invariants of the system, which need to hold as system progress. In Event-B these conditions are used to construct verification conditions for Invariant Preservation Rule. There are few of these rules, but the most important are Invariant Preservation Rule, Deadlock Freedom Rule and also satisfaction of refinement obligation. These set of rules just indicates what needs to be proven.

$$\left\{ \begin{array}{l} A(c) \\ I(c, v) \\ G(c, v) \end{array} \right\} \vdash I_i(c, E(c, v))$$

Fig. 3: Invariant Preservation Rule

In Figure 3, sequent of Invariant Preservation Rule is stated, where hypothesis consists of axioms, invariants and guards of the events. The goal is to validate whether invariants hold after event.

III. A CASE STUDY: ARMv6-M

Case study section intends to demonstrate potential of our domain specific language by analysing ARM Cortex M0+ microprocessor instruction set. This section starts with a brief description of the ARMv6-M instruction set and information we used to construct our model (Section III-A). Further, concrete example of translated instruction to is demonstrated together with visual graph representation (Sections III-B and III-C). As the basis for our experiment, we used microprocessor manual [11], which includes short descriptions of each instruction as well as functionality pseudocode and encodings. We used this information to build an instruction set model of this specific microprocessor.

A. Features & issues at the specification phase

The faithfulness of ARMv6-M model with the reference specification was our main concern. We aimed at representing each instruction listed in the ISA document using the framework presented in this work, with the least loss of information. This required to be as consistent as possible with the manual, replicating the flow of the operations with the same order, as

well as using the same names of the variables and functions at times.

As far as the reference documents are concerned, instructions characterised in there tend to come up with a bare description. Hence, sometimes being completely accurate with the source system is complex, also because of the lack of technical information about the original architecture. This was the case for some of the functions such as DataMemoryBarrier, BkptInstrDebugEvent, Hint_SendEvent, and so on [11]. None of them are described in the reference manual, therefore these functions are considered to be executed atomically (by one single node). Obviously, this might not be the case and a deeper research with the producers of the core is needed to untangle these operations.

Based on the structure described in Section II-A, we built the model on top of the microprogram module. Therefore, many of the instructions characterised are supported by the methods present at the bottom layers. For instance, the operation *incAndFetchInstruction* (depicted in the partial order in Figure 7) consists of three sequential micro-operations: incrementing the program counter (PC) register, writing the result back into such a register and eventually moving the content of the memory location pointed out by the PC into the instruction register (IR). This operation is executed at the end of nearly all the instructions, except for the ones involving any branching mechanisms. Another example which can be analysed is the *pop* operation. It is in charge of decrementing/incrementing the stack pointer (SP) register (according to the internal architecture), and then fetching from the memory location addressed by the SP a chunk of data. These two operations are used inside the main instruction set architecture.

Another part which is worth discussing is the one related to the variety of the registers present inside the core. Each of them has got a different function which we have attempted to replicate as much as we could. In this case, the purpose of each register is not described with much detail in the document. This partially limited the accuracy of the final model. Nonetheless, the high degree of abstraction we kept downsizes the importance of these problems. The registers present in our model are: linked register (LR), application program status register (APSR), interrupt program status register (IPSR), execution program status register (EPSR), the special-purpose mask register (also called PRIMASK), the special-purpose control register (CONTROL). They are needed to manage different internal functions. In our model, we limited at reading and writing them according to the reference description.

Additionally, all the arithmetic, logical and bitwise operations are represented in our model. The lack of technical details is present also in this side, limiting the accuracy of the model. We assumed that all these operations are performed via the arithmetic logic unit (ALU). Unfortunately, this is just an assumption, we are not sure whether the shift operation is truly accomplished through such a unit or via another dedicated one, for instance. This problem does not interfere with our main purpose of developing a systematic approach though. Indeed, the designer of a system is expected to have such a technical

knowledge, which can be employed to reach a high degree of faithfulness in the target model.

Given this general introduction about the model, in the next section some of the instructions characterised in our representation will be showed, pinpointing at the high level of similarity with the instruction set that we aimed at shaping.

B. Target model description

All the instructions included in the ARMv6-M have been characterised in our representation. According to the reference manual description, the ARMv6-M is a subset of ARMv7-M and supports Thumb 16-bit instruction set, including a small number of 32-bit instructions introduced to the architecture as part of the Thumb-2 technology in ARMv6T2.

Encoding T1

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
opcode7							imm3			Rn		Rd			
0	0	0	1	1	1	0	0	1	1	0	0	1	1	0	1

ADD R5, R1, #3

Encoding T2

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
opcode5					Rdn			imm8							
0	0	1	1	0	1	0	0	0	1	0	0	0	0	0	1

ADD R4, #65

Fig. 4: ADD (immediate) encodings - ARMv6-M ([11], p.107-108)

An interesting feature of the ISA is that some instructions have two opcodes, named T1 and T2. This provides a more flexible architecture, which in turn comes up with a flexible assembly language. This is the case for the *ADD (immediate)* instruction, for instance. The first encoding T1 has a seven bit opcode, and the rest of the instruction encoding is meant for hosting an immediate value (3-bit), and two more register locations (3-bit each): the source of the first operand and the destination of the result. The second encoding T2 instead comes up with a quite different pattern. The opcode associated to it is has only 5 bits, then there is a register (3-bit) which acts both as the source operand and the destination of the result; the remaining 8 bits contain a longer immediate value. Both of the encodings are depicted in Figure 4, where also an example of how these two encoding might be used is present. These instructions can be used for different purposes. Because of time constraints, we focused just on the instructions with the encoding T1. Nonetheless, we plan to represent the second encoding in future work, in order to create a fully realistic benchmark (see Section IV-A).

```

1 EncodingSpecificOperations();
2 (result, carry, overflow) = AddWithCarry(R[n], imm32, '0');
3 R[d] = result;
4 if setflags then

```

```

5 APSR.N = result<31>;
6 APSR.Z = IsZeroBit(result);
7 APSR.C = carry;
8 APSR.V = overflow;

```

Listing 1: ADD (immediate) instruction - Reference specification

The pseudocode specification of the instructions have been translated into Farfalle (described in Section II). During the conversion process, we assumed that all the parameters which are involved in the internal operations are available and ready to be used. This is the case for the *imm32* value, for instance, present in some of the instructions. This is over 32 bits, and is obtained extending the immediate value present inside the instruction, over 3, 5 or 8-bit long at times. We assumed to have a decoding function able to decode the instruction by the opcode. This is theoretically supposed to execute the right operations (such as the immediate extension) prior to the instruction execution.

Let us focus on the pseudocode specification of the ADD (immediate) instruction (showed in Listing 1). Three parameters are needed for this instruction: *R[n]*, *R[d]* and *imm32*. *setflags* is a flag always set to true for this instruction, therefore is not considered in our implementation. We can similarly characterise this instruction via the framework presented in this work, see Listing below.

```

1 add_ImmT1 rn rd imm32 = do
2   result <- alu (adcRegImm rn imm32 0)
3   writeRegister rd result
4   updateN result
5   updateZ result
6   updateC
7   updateV
8   incAndFetchInstruction

```

Listing 2: ADD (immediate) instruction - Farfalle specification

The instructions takes in input two registers and an immediate value over 32 bits. According to the reference document, the procedure *EncodingSpecificOperation()* is the one in charge of performing the right operations depending on the encoding read. As already mentioned, we assume this step to be already performed beforehand to instruction execution. All the other statements are quite similar to the source counterparts. The only thing worth mentioning is that the alu operation is in charge of setting the carry and the overflow bits internally. That is the reason why those two are not present in the Haskell implementation. The *IsZeroBit()* procedure simply checks if a value is 0, this is implemented inside the function *updateZ*. At the end of the Haskell specification is present the *incAndFetchInstruction*, which is in charge of updating the program counter register and fetching the next instruction.

Another instruction which might be of interest is the store, also known as STR (register). This instruction computes the

address where to store a chunk of data, from a base address and an offset, both defined into two registers. Although the offset can be optionally shifted in a general ISA architecture, this feature has not been implemented in the ARMv6-M. The *shift_n* variable in fact, which is supposed to contain the number of bits that the register $R[m]$ must be shifted by, is always set to 0. Below the source specification is showed:

```

1 EncodingSpecificOperations();
2 offset = Shift(R[m], shift_t, shift_n, APSR.C);
3 address = R[n] + offset;
4 MemU[address,4] = R[t];

```

Listing 3: STR (register) instruction - Reference specification

The specification based on Farfalle is listed in Listing 4. According to the source description, this instruction takes in input three registers. *rm* contains the base address, *rm* the optionally shifted offset (shifted by 0 bits in this implementation), and *rt* that contains the data chunk to be stored into the memory location pointed by the final result, eventually placed in the variable *address*. The arithmetic logic unit is in charge of performing either the shift and the addition, as described in Section III-A. *writeMemory* is one of the procedure implemented in the module basics (see Section II-A). *readBit* procedure, simply returns the bit 29, where the carry flag is stored.

```

1 str_RegT1 rm rn rt = do
2   offset <- alu (shlRegImm rm 0 (readBit apsr 29))
3   address <- alu (adcRegImm rn offset 0)
4   writeMemory address rt
5   incAndFetchInstruction

```

Listing 4: STR (register) instruction - Farfalle specification

The two Haskell-based specifications showed along this section are fairly similar to the reference pseudocode. The degree of precision the designer can go for depends on the accuracy the source specification have been characterised with. There has not been loss of information in the conversion process.

C. Visualisation

Going towards a more comprehensible ISA representation was one of our main goals. As discussed in Section I in fact, often instruction set architecture reference documents are hard to read and fully understand. We believe that having a visual representation of the operations which occur inside each instruction would make the understanding easier and quicker. That is the main reason why we chose to adopt partial orders to represent the behaviour of instructions. Each instruction can be analysed singularly. The behaviour can be isolated, focusing on the causality of the operations to execute.

The capability of representing sequential and concurrent behaviours is instrumental for modelling instructions. Partial orders address this requirement with the usage of nodes

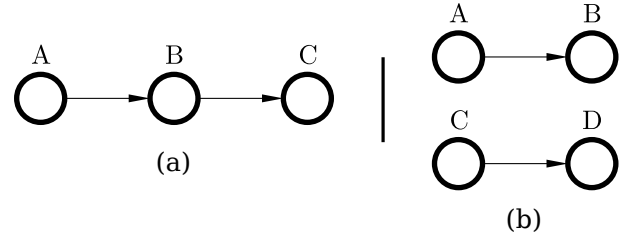


Fig. 5: Examples of sequential and concurrent behaviours.

(\circ) and dependencies (\rightarrow). Figure 5(a) depicts a partial order where the three operations are executed sequentially. Figure 5(b) instead shows four operations $\{A, B, C, D\}$ which can be executed with different orders. $A \rightarrow B$ and $C \rightarrow D$ are concurrent paths. A detailed description can be found in [2].

In the light of the above, we are working on automating the conversion process from the Farfalle specification to partial orders. The tool will be able to analyse the *ISA_module* file, capture the dependencies between the operations within each instruction and build a partial order.

Multiple levels of abstraction can be chosen for the conversion. Components such as the arithmetic logic unit can be adopted for the execution of different operations: i.e. addition, multiplication, shift, etc. One might want to represent each single operation separately with a different node (trading off some readability with accuracy), or raise the level of abstraction losing the internal information.

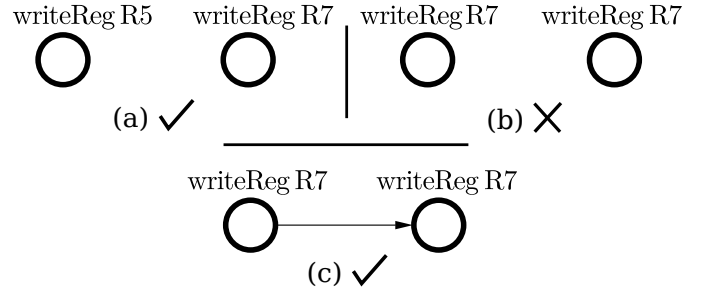


Fig. 6: Allowed and not allowed conversions.

The partial representation of the data, in some of the operations entailing registers and memory access, represents another layer of abstraction which can be traded for readability. This would allow the development of a more accurate model without overlapped register/memory accesses. For instance, let us assume that the register file has got two different input writing ports. This would allow two different registers to be written at the same time. The destination of a writing operation should be taken into consideration therefore, in order to avoid a simultaneous access at the same location of the register file (see Figure 6). The partial order in Figure 6(b) is not allowed because the two writing mechanisms can be triggered concurrently. A dependency should be inserted, as in the graph in Figure 6(c).

The following example clarifies the above issue. Listing 6 gives a specification of the AND (register) instruction (ARMv6-M ISA [11]).

```

1 and_RegT1 rm rdn = do
2   shifted <- alu (shlRegImm rm 0 apsr[29])
3   result <- alu (andRegImm rdn shifted)
4   writeRegister rdn result
5   updateN result[31]
6   updateZ result
7   updateC carry
8   incAndFetchInstruction

```

Listing 5: AND (register) instruction - Farfalle specification

This instruction can be automatically translated into the partial order in Figure 7. Every statement of the specification corresponds to a node in the graph. For instance, the one on fourth line correspond to the node named “writeReg”. The dependencies between the nodes are obtained by looking at the values which are needed for the execution of each operation. For instance, the second ALU operation needs the value called “shifted”, computed during the first ALU operation. Therefore, the second action depends on the first one. “incAndFetchInstruction” is a basic operation defined inside the module microprogram (see Section II-A). This is composed by three sequential operations, which are in this case extracted and depicted in the partial order, for an increased accuracy.

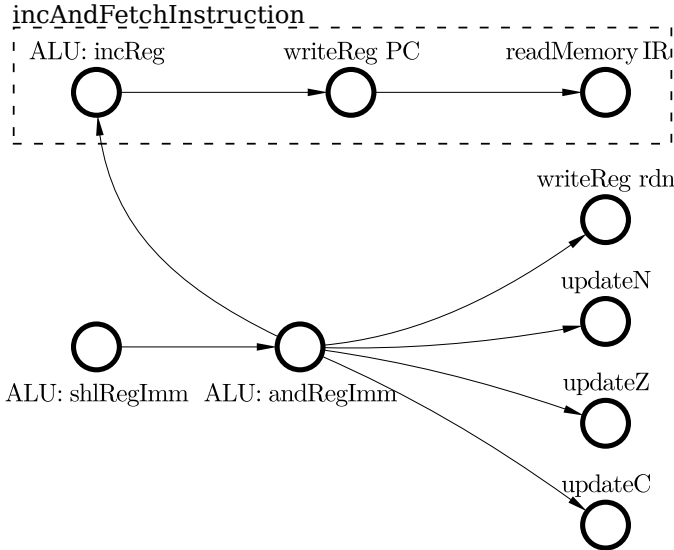


Fig. 7: AND (register) instruction - ARMv6-M ([11], p.116)

The graph in Figure 7 is more readable than the Haskell-based specification, listed in Listing 6. The dependencies between the various operations are clearly expressed via arcs. In addition, the graph-based model provides a good degree of fidelity; combined with the code-based specification it can lead to a good tradeoff between readability and accuracy.

Another example which worths a discussion is the unconditional branch. This instruction simply causes a branch to a target address. The Haskell-based specification is depicted below:

```

1 b imm32 = do
2   nextInstrLocation <- alu (adcRegImm pc imm32 0)
3   writeRegister pc nextInstrLocation
4   readMemory nextInstrLocation ir

```

Listing 6: B instruction - Farfalle specification

The particularity of this instruction is the lack of the *incAndFetchInstruction* step. The new instruction register indeed is obtained by the address computed internally, adding an 32-bit immediate address with the current program counter register. In Figure 8 the partial order associate to the branch instruction is showed. The reader should notice that there is no direct dependency between the *writeReg* and the *readMemory* nodes. Nonetheless, a dependency has been inserted in order to fetch the new instruction once that the PC register has been updated. This might be named a *functional dependency* and is implicit to the *incAndFetchInstruction* procedure either.



Fig. 8: Branch instruction - ARMv6-M ([11], p.119)

We did not have the time yet to implement the tool for the conversion. The example showed along this section has been obtained by hand. Therefore, we reserve for a future article the results and considerations related to the synthesis of the final controller.

D. Instruction Verification

This subsection briefly describes how instruction correctness can be validated. Constructing Event-B model is refinement driven process, where developer starts with the most abstract concept of the system. We started by declaring very basic concepts of microprocessors system. A typical, Event-B model structure consists of context and machine frames, our initial context was used to declare some trivial constants and data types e.g. microprocessors’ word size. Original context then was refined by adding more information about operation type structure, decoding opcode/operands functions etc. For instance, in (1), we state that instruction is pair type with opcode and three operands, whereas i.e. opcode-decode function can be defined as shown in (2).

$$(opcode, (value, (value, value))) \in OPERATION \quad (1)$$

$$DECODE \in OPERATION \rightarrow OPCODE \quad (2)$$

After declaration of variables, types and structures in the context file, machine file was created, which extended those

concepts as well as included events. Machine model required further expansion of some key concepts of data, registers, program counter, but most importantly it contained instruction events. In Event-B, event can be conditional or unconditionally fired, depending on system requirements for that event. In the context of the microprocessor and our model, event of the particular instruction is executed if opcode satisfies the guard of the event. For example, (3) shows one of the guards of the ADD instruction.

$$DECODE(CODE(pc)) = ADD \quad (3)$$

Satisfying guard condition means action in the event can be executed. In the context of ADD instruction simple addition of two source registers is performed and the sum is loaded to the destination register, which is then followed by increment of program counter register.

However, analysis of ARM Cortex M0+ instruction set showed us that even simple instructions like addition are difficult to specify formally in Event-B. Therefore, to make our model more feasible another refinement step would be necessary in order to capture the same level of details as in specifications. Another important requirement for instruction validation would be translation automation. Currently, we had to define the system and correctness conditions as well as prove them. As stated before it is our objective to hide this mathematical modelling exercise from engineers. Hence, together with creating a general framework, we need to automate this process.

IV. CONCLUSION

This work represents the first step towards a formal and readable specification of processor instruction sets. We demonstrated how the specification of a real ISA, which has been embedded into the ARM Cortex-M0, can be specified via our framework without losing any information or lowering the accuracy of the representation.

In addition, we showed how such a framework can be converted into a visual-based representation with different tradeoffs for maximising either the readability or the accuracy of the graph-based model. The latter, based on the Conditional Partial Order Graph, can exploit a built-in tool-chain embedded in *Workcraft* for visualisation and hardware synthesis purposes. In particular, this representation aims at simplifying the understanding of the processor instruction sets.

This open-source framework [17] opens the path towards different other languages and tools which can be adopted in the design flow. Event-B might be one such piece of the puzzle which we hope to be able to complete, together with the whole community in this field.

A. Future research directions

Much research can be carried on about the work presented. First of all we will expand the framework developed, providing some procedures to capture the internal instruction dependencies, automatically deriving a partial order according to the level of abstraction chosen by the user. This will be used

to move from the DSL - Haskell framework, into the CPOG representation (see Figure 2), avoiding likely human errors. In turn, this will allow us to derive the hardware microcontroller and make a consistent comparison (in terms of area, power consumption, latency, etc.) with the one which the ARM Cortex-M0 embeds.

Secondly, we will work on a plugin to automatically convert the specification into the Event-B language. This will allow the usage of sound theorem provers for the formal verification of the ISA, an extremely useful instrument which designers might benefit from.

From the point of view of the case study itself, the model will be expanded to include the T2 encodings. Other case studies will be carried out, such as the instruction sets of the Texas Instruments MSP430 and the Intel 8051 processors. As is pointed out in [14], nowadays there are quite a few architectures with a relevant market size - by ARM, ARC, Xtensa, MIPS. Unit volumes are progressively increasing, hence we believe the need for modelling and experimenting by these companies will do so. Eventually, designers and end users will benefit from a systematic approach to the design of processor instruction sets.

REFERENCES

- [1] A. Mokhov, A. Yakovlev. "Conditional partial order graphs: Model, synthesis, and application". IEEE Transactions on Computers, Volume 59, Pages 1480-1493, November 2010.
- [2] A. Mokhov. "Conditional Partial Order Graphs". Ph.D. Thesis, Newcastle University, September 2009.
- [3] I. Poliakov, D. Sokolov, A. Mokhov. "Workcraft: A static data flow structure editing, visualisation and analysis tool". Petri Nets and Other Models of Concurrency - ICATPN 2007. Pages 505-514, 2007.
- [4] A. Mokhov, A. Alekseyev, A. Yakovlev. "Encoding of processor instruction sets with explicit concurrency control". Computers & Digital Techniques, IET. Volume 5, Pages 427-439. November 2011.
- [5] A. de Gennaro, P. Stankaitis, A. Mokhov. "A Heuristic Algorithm for Deriving Compact Models of Processor Instruction Sets". 15th International Conference on Application of Concurrency to System Design (In Press). June 24-26, 2015.
- [6] A. Fox, M. Myreen. "A Trustworthy Monadic Formalization of the ARMv7 Instruction Set Architecture". Interactive Theorem Proving (ITP), pages 243-258, 2010.
- [7] F. Yuan, K. Eder. "A Generic Instruction Set Architecture Model in Event-B for Early Design Space Exploration". Technical Report CSTR-09-006, University of Bristol, September 2009.
- [8] T. A. Cook, E. Harcourt. "A Functional Specification Language for Instruction Set Architectures". Proceedings of the 1994 International Conference on Computer Languages. Publisher IEEE, pages 11-19. May 1994.
- [9] A. Abbas, A. Ahmed, A. Ahmed, W. Uz Zaman Bajwa, A. Anwar, S. Abbasi. "A retargetable tool-suite for the design of application specific instruction set processors using a machine description language". IEEE International Symposium on Circuits and Systems, 2002. Volume 1, pages 425-428. ISCAS 2002.
- [10] M. Rykunov. "Design of Asynchronous Microprocessor for Power Proportionality". Ph.D. Thesis, Newcastle University. Technical Report Series NCL-EEE-MICRO-TR-2013-182, December 2013.
- [11] ARM Ltd. "ARMv6-M Architecture Reference Manual". ARM DDI 0419C (ID092410), 2010.
- [12] S. P. Jones (editor) et al. "Haskell 98 Language and Libraries The Revised Report". December 2002.
- [13] C. Mtayer, J.-R. Abrial, L. Voisin. "Event-B Language". Project IST-511599. 31st May 2005

- [14] J. Turley. “*CPUs Get Cheaper; Transistors Get More Expensive*”. Electronic Engineering Journal: <http://www.eejournal.com/archives/articles/20151104-synopsys>. November 4, 2015
- [15] Abrial, Jean-Raymond. “*Modeling in Event-B: System and Software Engineering*”. Cambridge University Press, New York, NY, USA. 2010
- [16] J. Welte, H. Manz. “*openETCS D2.1: Report on existing methodologies*”. ITEA2 Project Technical Report. January 2012
- [17] A. Mokhov, A. de Gennaro, P. Stankaitis. <https://github.com/allegroCoder/Farfalle>. Github repository.